

Jeffrey Zhang

San Diego, CA | (858) 231-6881 | jeffreyzhangsd@gmail.com | [LinkedIn](#) | [Github](#) | [Portfolio](#)

Education & Certifications

University of California, Los Angeles

B.S. Applied Mathematics

Oct 2020 - Jun 2022

Hack Reactor

Certificate in Advanced Software Engineering

Jan 2023

Technical Skills

Front End: React, JavaScript, TypeScript, Next.js, React Native, Swift, Kotlin, Tailwind, Redux, HTML/CSS, SEO

Back End: Python, Java, C++, Node.js, Express, PostgreSQL, MongoDB, SQL, MySQL, REST APIs, Firebase

Testing/Deployment: Docker, AWS EC2, TDD, Jest, Mocha, Chai

Tools & Technologies: Vim, Git, Webpack, Babel, Nginx, Agile, Jira

Experience

Software Engineer | ANSOL Inc.

Oct 2024 - Present

- Migrated a legacy planning and scheduling application to a modern JavaScript-based framework, improving front-end performance, maintainability, and data integration across key system modules.
- Designed and delivered new user-facing features including dynamic timelines, sortable tables, and modals, directly improving analyst productivity and streamlining complex data workflows.
- Guided an external team's security compliance web application and spearheaded its migration through view-model design, data typing, and component integration, significantly accelerating their migration timeline and eventually becoming the point of contact for the app.

Software Engineer Intern | Checkins.ai

Sep 2023 - Mar 2024

- Achieved a 5x boost in website speeds by migrating codebase to NextJS, leveraging server-side rendering to drastically improve initial page loading, interactivity and overall user experience.
- Developed a scalable design system using Tailwind, ensuring consistency and responsiveness across various devices for a 25% reduction in design inconsistencies, along with SEO best practices.

Projects

Full Stack Software Engineer | Local AI Agent Platform

Apr 2026 - Present

Self-hosted multi-agent automation hub running tiered local LLMs on a 16GB Mac Mini.

- Architected a three-tier local model router (Ollama + llama.cpp) that triages email and messages, grades job postings, and compresses research locally, cutting cloud LLM spend by 30-40%.
- Served a 35B-parameter model on 16GB RAM via mmap SSD streaming with zero swap (~17 tok/s on Apple M4), enabling daily long-context summarization on consumer hardware.
- Built a fleet of scheduled daemons — SEC-filing trading signals with paper-trading execution, knowledge-base ingestion, and self-auditing health reports — coordinated through lock-based scheduling.

Front End Software Developer | Sip & Play Refreshed | [GitHub Repo](#) | [Link to Codédex](#) | [Demo](#)

Jul 2024

Second place winner for the Codédex Summer Hackathon with focus on 3D models, approachable UI/UX and responsiveness.

Back End Software Engineer | Gaviali | [GitHub Repo](#)

Jan 2023

A scalable back end with routes and database capable of managing 1000+ active users.

- Initialized a PostgreSQL database with Express routes on Amazon EC2 instances to handle millions of rows of data and handle a rate of 1000 API requests per second and more.
- Reduced response times from 6s to <200ms with optimized queries and clustering
- Deployed a Nginx load balancer with caching across Express servers for efficient PostgreSQL queries.